

Tessera Normative Orientation Architecture

Development and Holistic Measurement Specification

Document ID: TNOA-DMS-0.1

Version: 0.1 Provisional

Date: 2026-07-11

Status: Active / Provisional / Expansion Scaffold

Purpose

Define the staged development, independent construct analysis, semantic representation, measurement strategy, and validation plan for the Tessera Normative Orientation Architecture.

Contents Overview

1. Executive summary and scientific position
2. Governing construct, scope, and layered architecture
3. Bioecological organization
4. Independent construct workstreams
5. Cross-cutting axes and Construct Development Packet
6. Normative Decision Episode and graph implementation
7. Measurement battery and Tessera integration
8. Scoring policy, validation roadmap, and profile output
9. Repository implementation, risks, and appendices

Overwrite and reinterpretation notice

This specification introduces **Normative Orientation Architecture** as the descriptive umbrella for values, principles, norms, ideology, morality, ethics, factual premises, motives, intentions, behavior, and developmental updating.

It does **not** delete the existing Karma model or personality taxonomy. It repositions them within a clearer architecture:

- Personality describes dispositional tendencies and characteristic forms of expression.
- TNOA describes how an individual constructs and applies judgments of requirement, permission, prohibition, justification, blame, entitlement, sacrifice, and responsibility.
- Karma remains a provisional ethical-review and evaluative overlay rather than the primary descriptive representation of normative identity.
- Reconciliation remains a mechanism for examining mismatches, ambiguity, competing explanations, and context without assuming error or moral failure.

Affected repository areas include the project overview, model manifest, Karma boundary language, specification index, and documentation workflows. Existing source material is preserved and remains reviewable.

1. Executive summary

TNOA is intended to answer:

How does a person construct, prioritize, justify, and enact judgments about what is required, permitted, prohibited, excusable, praiseworthy, blameworthy, or desirable across contexts and over time?

The architecture must not begin by deciding whether a person is good or bad. Its first task is to reconstruct how conduct becomes justifiable to the individual. Ethical evaluation may be applied later through an explicitly declared benchmark.

The model therefore separates three layers:

- 1 **Descriptive layer:** What the person believes, values, permits, prohibits, intends, and excludes.
- 2 **Explanatory layer:** How personality, cognition, affect, development, relationships, learning, ideology, and ecology contribute to those orientations.
- 3 **Evaluative layer:** How the orientation compares with a stated ethical, legal, professional, cultural, or human-rights standard.

The core unit of analysis is the **Normative Decision Episode**, not a global morality score. A complete profile is assembled from repeated episodes across roles, relationships, audiences, ecological systems, levels of constraint, and time.

TNOA should be regarded as an integrative developmental identity architecture and hypothesis-generating framework until its constructs, measures, processes, and interpretations are empirically validated.

2. Scientific position and evidence governance

2.1 Evidence levels

TNOA adopts the Tessera evidence-classification approach.

Level	Classification	Meaning for TNOA
1	Empirical foundation	Established research domains may support individual variables or processes, such as values, norms, attitudes, moral judgment, intention, social influence, learning, memory, stress, and developmental context.
2	Integrative framework	TNOA combines supported constructs into an organized decision architecture. The integration may be theoretically justified without yet being validated as one unified model.
3	Tessera hypothesis	New graph structures, composite variables, automated inference, profile synthesis, contextual coefficients, and developmental update rules require direct testing.

2.2 Evidence rules

- Empirical findings must not be overwritten merely to simplify the model.
- A construct supported in one research context must not be assumed valid in every culture, age group, or application.
- Source-informed does not mean psychometrically validated as a Tessera measure.
- Machine-readable does not mean inferentially reliable.
- The model must permit unknown, not assessed, not applicable, and insufficient evidence states.
- Every inference should preserve its source, date, context, confidence, and alternative explanations.
- No single observation should be converted into a stable identity claim without sufficient repeated evidence.

2.3 Source lineage

This specification develops the project direction established by:

- Tessera Technical Specification v1.0.
- Tessera Developmental Identity Framework: Evidence Classification and Cross-Validation Reference.
- Integrated Psychodevelopmental Life Model.
- Resources for Developmental Ecological Psychology.
- Ethical Integrity 12-Dimensional Matrix Model.
- Repository governance, personality, ecology, identity, Karma, and reconciliation documents.

The architecture also draws conceptually from bioecological, relational-developmental, sociocultural, social-cognitive, narrative-identity, moral-psychology, decision, motivation, intention, and behavioral-learning traditions. These source traditions remain distinct from the Tessera-specific synthesis.

3. Governing construct

3.1 Formal definition

Normative orientation is the patterned way an individual assigns authority, priority, permission, obligation, prohibition, entitlement, praise, blame, excuse, and responsibility to possible actions, relationships, outcomes, institutions, identities, and social arrangements.

Normative Orientation Architecture is the organized system of values, principles, norms, ideologies, factual premises, moralized concerns, ethical reasoning modes, motives, intentions, contextual conditions, and developmental processes through which those assignments are produced and revised.

3.2 Foundational semantic distinction

The architecture distinguishes:

- **Values:** Desired ends, conditions, relationships, identities, or states of affairs.
- **Principles:** Generalized rules that translate values into judgments.
- **Norms:** Perceived standards of expected, approved, or role-appropriate conduct.
- **Ideology:** A larger worldview connecting social categories, history, authority, causality, identity, threat, legitimacy, and prescriptions for social order.
- **Morality:** Judgments experienced as possessing special force involving obligation, prohibition, entitlement, praise, blame, or deserved sanction.
- **Ethics:** Reflective or systematized organization of normative judgments and reasons.
- **Factual premises:** Beliefs about what is true, likely, causal, threatening, beneficial, or historically relevant.
- **Motives:** Desired outcomes or states energizing action.
- **Intentions:** Selected commitments to attempt particular actions.
- **Behavior:** Enacted or omitted conduct under actual opportunities and constraints.
- **Consequences:** Intended, foreseen, unforeseen, perceived, and actual outcomes.

3.3 What morality means in this architecture

Morality is not equated with compassion, harmlessness, conformity, productivity, legality, or social approval. These may become moral contents, but none defines morality by itself.

At its most foundational, morality is a system through which an agent gives certain reasons special normative force over conduct and treatment. An issue becomes moralized when it is experienced not merely as liked, useful, customary, or legal, but as required, forbidden, deserved, blameworthy, praiseworthy, or universally binding.

A coherent, internalized moral system may support humane conduct or harmful ideology. Internalization, conviction, complexity, consistency, and cost-bearing commitment are descriptive properties, not proof of ethical soundness.

4. Scope and non-goals

4.1 In scope

TNOA will represent:

- normative content;
- judgment structure;
- context-sensitive activation;
- developmental provenance;
- public-private divergence;
- intention-behavior divergence;

- factual-premise dependence;
- responsibility and excuse models;
- harm and beneficiary structure;
- revisability and identity protection;
- ecological and historical change;
- evidence confidence and uncertainty.

4.2 Out of scope for initial development

TNOA will not initially provide:

- a global good-versus-evil score;
- diagnosis of moral character;
- proof of hidden intent;
- automatic conversion from personality traits to morality;
- automatic criminal, clinical, employment, or legal conclusions;
- universal rankings of cultures or ideologies;
- deterministic prediction of behavior;
- claims that articulate reasoning necessarily produces ethical action;
- claims that social conformity demonstrates internal moral endorsement.

5. Layered architecture

5.1 Level A: Normative source systems

These provide the content from which judgments are constructed:

- values;
- principles;
- norms;
- ideologies and worldviews;
- moralized concerns;
- ethical frameworks;
- identity commitments;
- factual and causal beliefs;
- role expectations;
- narratives of history, threat, legitimacy, and belonging.

5.2 Level B: Decision-forming processes

These determine what becomes active in an episode:

- perception and situation construction;

- schema and association activation;
- affective appraisal;
- value activation;
- principle selection;
- norm recognition;
- authority recognition;
- factual-premise retrieval;
- outcome forecasting;
- beneficiary and harm representation;
- responsibility attribution;
- justification;
- deontic classification;
- motive competition.

5.3 Level C: Action formation

These determine whether orientation becomes conduct:

- intention strength;
- implementation planning;
- perceived behavioral control;
- capability;
- opportunity;
- role power;
- constraint load;
- habit and behavioral momentum;
- audience and social pressure;
- risk and sanction expectations;
- actual execution or omission.

5.4 Level D: Developmental updating

These explain how an episode changes the system:

- consequence evaluation;
- reinforcement or punishment;
- guilt, shame, pride, resentment, relief, fear, or moral elevation;
- social approval, rejection, or exclusion;
- narrative revision;
- identity consolidation or threat;
- value reprioritization;

- ideological reinforcement or disconfirmation;
- principle modification;
- accessibility shift;
- associative reweighting;
- constraint formation or relaxation.

5.5 Core processing sequence

Person-in-context

- > Perception and situation construction
- > Factual and causal interpretation
- > Value, norm, ideology, and identity activation
- > Principle or justificatory mode selection
- > Deontic judgment
- > Motive competition
- > Intention and plan
- > Behavior under capability, opportunity, and constraint
- > Consequence
- > Appraisal, narrative integration, and developmental update

The architecture treats the person and context as a dynamic unit. Context is not a decorative variable attached after a judgment has already been scored.

6. Bioecological organization

TNOA uses the Process-Person-Context-Time model as a developmental organizing frame.

6.1 Process

Proximal processes are recurring interactions through which normative structures become practiced and internalized. Examples include:

- caregiver discipline and explanation;
- conflict and repair;
- peer inclusion and exclusion;
- role modeling;
- religious ritual;
- school discipline;
- workplace incentives;
- political discussion;
- exposure to media narratives;
- participation in institutions;
- repeated encounters with authority, scarcity, privilege, or threat.

6.2 Person

Person characteristics include:

- temperament and personality;

- cognitive and linguistic capacity;
- attachment expectations;
- affective responsiveness;
- executive control;
- social identity;
- prior learning;
- health and stress state;
- perceived agency;
- capabilities and vulnerabilities.

These characteristics affect exposure, interpretation, elicited responses from others, and the resources available for reflection. They do not determine normative content by themselves.

6.3 Context

Microsystem

Immediate face-to-face settings such as family, peers, school, religious community, workplace, intimate relationships, and local groups.

Mesosystem

Relations among microsystems. Contradictory signals between family, school, workplace, peer group, and community may produce compartmentalization, conflict, or selective rule application.

Exosystem

Settings that affect the individual indirectly, including employment conditions, media systems, institutional policy, neighborhood resources, legal structures, and the decisions of powerful organizations.

Macrosystem

Cultural narratives, ideologies, economic arrangements, social hierarchies, religious cosmologies, law, historical memory, and accepted models of personhood and authority.

Chronosystem

Developmental timing, cohort, historical events, transitions, political change, technological change, migration, trauma, conversion, and cumulative exposure.

6.4 Time

The model distinguishes:

- microtime: what happens within an episode;
- mesotime: how patterns repeat across days, months, or recurring roles;
- macrotime: developmental, cohort, and historical change.

A complete profile must therefore include both stable tendencies and within-person variation.

7. Construct development program

Every umbrella element must be developed as an independent workstream before holistic scoring or automated synthesis is attempted.

7.1 Workstream 1: Values

Core question

What ends, states, relationships, identities, or conditions does the person regard as important?

Required distinctions

- abstract values versus concrete goals;
- desired values versus enacted values;
- personal versus collective values;
- intrinsic versus instrumental values;
- stable importance versus momentary salience;
- declared importance versus trade-off priority;
- identity-central versus peripheral values;
- protected versus negotiable values.

Measurement candidates

- open value elicitation;
- ranking and best-worst scaling;
- forced trade-offs;
- resource-allocation tasks;
- narrative examples;
- ecological momentary sampling;
- value-conflict scenarios;
- beneficiary substitution.

Output

A **value hierarchy network** containing centrality, contextual activation, conflicts, compatibility, beneficiaries, sacrifice thresholds, and developmental change.

7.2 Workstream 2: Principles

Core question

What generalized rules does the person use to translate values into judgments?

Required distinctions

- absolute versus defeasible principles;
- universal versus role-specific principles;
- positive duties versus prohibitions;

- procedural versus substantive principles;
- publicly stated versus privately applied principles;
- genuine rules versus post hoc rationalization;
- self-directed versus other-directed standards.

Measurement candidates

- rule elicitation;
- matched scenarios;
- exception probing;
- counterexample testing;
- reversal tests;
- self-other symmetry tests;
- in-group/out-group substitution;
- consistency across domains.

Output

A **principle application map** with trigger conditions, exceptions, priority, intended scope, actual scope, and conflict-resolution rules.

7.3 Workstream 3: Norms

Core question

What standards does the person perceive as socially expected, approved, or role appropriate?

Required distinctions

- descriptive norms: what people usually do;
- injunctive norms: what people approve or demand;
- personal norms: what the individual believes they ought to do;
- role norms: what is expected from a parent, citizen, leader, employee, believer, or group member;
- institutional norms: what organizations reward or punish;
- compliance versus endorsement;
- resistance versus concealment.

Measurement candidates

- perceived norm estimates;
- approval estimates;
- sanction expectation;
- anonymous versus identified response;
- audience manipulation;
- authority observation;
- dissent availability;

- private agreement and desire-for-change probes.

Output

A **normative pressure profile** showing internal endorsement, strategic compliance, fear-based compliance, role conformity, resistance, and public-private divergence.

7.4 Workstream 4: Ideology and worldview

Core question

What larger system connects beliefs about reality, identity, history, authority, groups, causes, and preferred social order?

Required components

- social categories;
- group identities;
- historical narratives;
- causal explanations;
- threat models;
- legitimacy claims;
- authority structures;
- value hierarchies;
- sacred commitments;
- prescriptions for social organization;
- narratives of enemies, victims, heroes, and decline;
- models of human nature.

Required distinctions

- content versus strength;
- coherence versus factual accuracy;
- identity commitment versus reflective endorsement;
- institutional ideology versus personal adaptation;
- symbolic affiliation versus operational belief;
- ideology versus isolated prejudice;
- ideology versus religion, while allowing overlap.

Output

An **ideological meaning network** connecting beliefs, values, identities, authorities, threats, prescriptions, and exception rules.

7.5 Workstream 5: Morality and moralization

Core question

Which issues does the person experience as possessing special normative force?

Required distinctions

- preference versus moral conviction;
- convention versus morality;
- prudential concern versus moral prohibition;
- sacred concern versus interpersonal concern;
- legal prohibition versus moral condemnation;
- disgust versus wrongdoing;
- identity threat versus moral violation;
- strong attitude versus perceived universal obligation.

Measurement candidates

- wrong-versus-undesirable probes;
- authority-removal tests;
- consensus-removal tests;
- legality-change tests;
- harm-removal tests;
- punishment and tolerance judgments;
- perceived universality;
- willingness to impose the rule on others;
- willingness to bear cost.

Output

A **moralization map** showing domains, strength, perceived objectivity, sanction support, universality, and reversibility.

7.6 Workstream 6: Ethics and justificatory repertoire

Core question

What consciously articulated modes does the person use to organize or defend normative judgments?

Candidate modes

- consequence-oriented;
- rule-oriented;
- rights-oriented;
- virtue and character-oriented;
- care and relationship-oriented;
- contractual and reciprocity-oriented;
- sacred or divine-command-oriented;
- role and professional ethics;
- precedent and tradition;

- necessity or emergency reasoning;
- particularistic case reasoning.

Measurement candidates

- dilemma interviews;
- justification coding;
- rule-consequence conflicts;
- relationship-universality conflicts;
- role comparisons;
- think-aloud protocols;
- confidence and revisability ratings;
- alternative-justification generation.

Output

A **justificatory repertoire profile**, not a single stage or philosophical type.

7.7 Workstream 7: Factual and epistemic premises

Core question

What does the person believe is true about the situation, actors, risks, groups, causes, and likely outcomes?

Required distinctions

- observed versus inferred belief;
- testimonial versus directly experienced belief;
- stereotype versus individuated evidence;
- probability estimate versus categorical claim;
- causal assumption versus association;
- historical claim;
- metaphysical or religious claim;
- confidence versus accuracy;
- evidence standard;
- source trust;
- willingness to update.

Measurement candidates

- premise elicitation;
- source tracing;
- confidence judgments;
- calibration tasks;
- evidence-update tasks;

- counterevidence response;
- causal diagramming;
- ambiguity manipulation.

Output

An **epistemic premise graph** attached to each judgment. The model may describe coherence relative to premises without endorsing those premises as accurate.

7.8 Workstream 8: Motives

Core question

What immediate or intermediate outcome does the person want?

Candidate motives

- protection;
- affiliation;
- status;
- control;
- revenge;
- care;
- justice;
- order;
- purity;
- security;
- approval;
- shame avoidance;
- material gain;
- group advancement;
- identity confirmation;
- meaning preservation;
- autonomy;
- relief from distress.

Required distinctions

- conscious versus nonconscious;
- proximal versus distal;
- self-reported versus inferred;
- primary versus secondary gain;
- mixed motives;
- stable motive tendency versus episode-specific motive;

- moral motive versus strategic motive.

Output

A **motive competition profile** showing activation, priority, inhibition, conflict, and retrospective reconstruction.

7.9 Workstream 9: Intentions

Core question

What action has the person committed to attempting?

Required distinctions

- desire versus intention;
- inclination versus commitment;
- abstract intention versus implementation plan;
- immediate versus delayed intention;
- conditional versus unconditional intention;
- private versus declared intention;
- prospective report versus retrospective claim;
- intention strength versus perceived control.

Measurement candidates

- action commitment ratings;
- implementation-intention prompts;
- time horizon;
- means and contingency planning;
- perceived control;
- opportunity manipulation;
- follow-up behavior tracking.

Output

An **intention specification** containing intended act, target, beneficiary, outcome, foreseen side effects, accepted costs, timing, perceived control, and commitment strength.

7.10 Workstream 10: Behavior and consequences

Core question

What was enacted or omitted, under which opportunities and constraints, and what followed?

Required distinctions

- attempted versus completed behavior;
- voluntary versus coerced behavior;
- habitual versus deliberative behavior;
- public performance versus private enactment;

- commission versus omission;
- capability failure versus intention failure;
- immediate versus delayed consequence;
- intended versus unintended outcome;
- perceived versus actual outcome.

Output

An **enactment profile** linking orientation, intention, control, opportunity, behavior, consequence, and later revision.

8. Cross-cutting descriptive axes

The following dimensions are coding axes applied to episodes. They are not assumed to be fixed personality traits.

Axis	Central question
Normative scope	Who or what receives consideration or standing?
Concern priority	Which values, interests, identities, or institutions dominate?
Deontic assignment	Is the act required, permitted, optional, discouraged, or forbidden?
Authority locus	What grants the judgment legitimacy?
Justificatory form	What type of reason supports the judgment?
Epistemic premise structure	What must be believed for the judgment to make sense?
Beneficiary orientation	Who is intended to gain or be protected?
Harm and means tolerance	What costs, coercion, or harms are accepted, and toward whom?
Responsibility model	How are control, intention, foreseeability, excuse, blame, praise, and repair assigned?
Internalization and conviction	Is the judgment externally imposed, personally endorsed, or identity-central?
Contextual modulation	Which changes in role, audience, target, power, threat, scarcity, or institution alter it?
Revisability	What evidence, experience, or reflection could change it?

These axes should support continuous, categorical, network, and qualitative representations where appropriate. They should not be summed into a global morality score by default.

9. Standard Construct Development Packet

Every workstream must complete the same packet before integration.

9.1 Conceptual specification

- formal definition;
- plain-language definition;
- necessary features;
- optional features;

- exclusion criteria;
- neighboring constructs;
- common conflation;
- classification as state, trait, process, structure, outcome, or mixed construct.

9.2 Facet structure

- proposed facets;
- reflective versus formative status;
- expected relations among facets;
- conditions under which facets dissociate;
- cultural and developmental variation;
- candidate higher-order relations.

9.3 Developmental account

- earliest precursors;
- typical developmental transformations;
- sensitive periods or developmental dependencies;
- socialization pathways;
- identity integration;
- adult revision;
- possible aging-related changes;
- equifinality and multifinality.

9.4 Bioecological account

- proximal processes;
- person characteristics;
- microsystem sources;
- mesosystem conflicts;
- exosystem pressures;
- macrosystem narratives;
- chronosystem changes;
- repeated person-context transactions.

9.5 Mechanistic account

- automatic activation;
- reflective processing;
- affect;
- memory;

- learning and reinforcement;
- identity protection;
- constraint load;
- accessibility;
- capability and control;
- behavior-consequence feedback.

9.6 Measurement specification

- self-report;
- interview;
- scenario task;
- behavioral task;
- informant data;
- naturalistic observation;
- longitudinal sampling;
- process measures;
- scoring and coding rules;
- missing-data handling;
- uncertainty representation.

9.7 Validity risks

- social desirability;
- demand characteristics;
- cultural bias;
- literacy and education effects;
- ideological framing;
- retrospective distortion;
- impression management;
- confusion with intelligence;
- confusion with conformity;
- inference beyond evidence;
- evaluator bias;
- criterion contamination.

9.8 Tessera implementation

- object types;
- link types;

- accessibility variables;
- constraint variables;
- temporal properties;
- source provenance;
- confidence level;
- evidence classification;
- graph queries;
- reporting rules.

9.9 Advancement gate

A construct may advance to integration only when:

- 1 its definition and boundaries are approved;
- 2 its closest neighboring constructs are explicitly differentiated;
- 3 its measurement modes are matched to its structure;
- 4 foreseeable bias risks are documented;
- 5 the evidence level is declared;
- 6 the graph representation does not imply more certainty than the evidence supports.

10. Normative Decision Episode

The primary semantic unit is an episode rather than a person-level score.

10.1 Human-readable template

In context C, toward target T, the person regards action A as required, permitted, optional, discouraged, or forbidden because of reason R, grounded in value V and authority S, given factual premise B, intending outcome O, while accepting or rejecting consequence H, with confidence Q and cost-bearing willingness K.

10.2 Required fields

episode_id
actor
observer_or_source
perceived_situation
objective_context_if_available
role
audience
target
target_identity_or_category
relationship
power_differential
available_actions
selected_action
normative_domain
normative_status
activated_values
applied_principles
perceived_descriptive_norms
perceived_injunctive_norms
personal_norms
ideological_frame
authority_source
factual_premises
causal_premises
expected_outcomes
immediate_motives
intended_beneficiary
intended_burden
foreseen_harms
accepted_harms
rejected_alternatives
responsibility_judgment
intention_strength
implementation_plan
perceived_control
capabilities
opportunities
constraint_load
actual_behavior
actual_outcome
perceived_outcome
post_event_justification
emotional_response
narrative_meaning
identity_implication
revision_or_reinforcement
evidence_source
confidence
timestamp
validation_status

10.3 Semantic safeguards

- intended_burden must remain separate from foreseen_harm.
- perceived_outcome must remain separate from actual_outcome.
- post_event_justification must not overwrite the prospective intention record.
- public_response and private_response must remain separately representable.
- inferred motives or intentions must be labeled as inference, not direct report.
- target-group labels should preserve the participant's categorization and a neutral research coding, where appropriate.

- harmful or false beliefs may be represented descriptively without being normalized or endorsed.

11. Graph implementation

11.1 Core nodes

- Person
- Role
- Relationship
- Group
- Institution
- Context
- Event
- Action
- Target
- Value
- Principle
- Norm
- Ideology
- Moralized Concern
- Ethical Mode
- Factual Premise
- Causal Premise
- Motive
- Intention
- Capability
- Constraint
- Behavior
- Consequence
- Emotion
- Narrative
- Identity Commitment
- Evidence Source

11.2 Core relations

Person -> endorses -> Value
Person -> applies -> Principle
Person -> perceives -> Norm
Person -> identifies_with -> Ideology
Person -> believes -> Factual Premise
Person -> occupies -> Role
Person -> relates_to -> Target

Judgment -> concerns -> Action
Judgment -> targets -> Person/Group/Object
Judgment -> assigns -> Normative Status
Judgment -> justified_by -> Reason
Judgment -> authorized_by -> Authority Source
Judgment -> activated_in -> Context
Judgment -> prioritizes -> Value
Judgment -> tolerates -> Harm
Judgment -> attributes -> Responsibility

Judgment -> generates -> Intention
Intention -> implemented_by -> Plan
Intention -> constrained_by -> Constraint
Intention -> enabled_by -> Capability
Intention -> expressed_as -> Behavior
Behavior -> produces -> Consequence
Consequence -> updates -> Narrative/Association/Value/Principle

11.3 Link metadata

Every relation may include:

- strength;
- confidence;
- valence;
- source;
- first-observed and last-observed dates;
- context specificity;
- frequency;
- recency;
- contradiction markers;
- validation status;
- privacy classification.

11.4 Computational restraint

Graph centrality, clustering, similarity, and prediction may assist exploration. They must not be presented as validated psychological meaning until tested against independent criteria and human review.

12. Measurement battery

No single questionnaire should be treated as sufficient.

12.1 Instrument 1: Normative orientation interview

Purpose:

- elicit values, rules, loyalties, prohibitions, authorities, and exceptions in the participant's own language;
- preserve emic meaning before standardized coding;
- identify personally salient domains;
- generate candidate scenarios for later testing.

Primary outputs:

- self-defined values;
- principle statements;
- authority sources;
- moralized domains;
- exception rules;
- identity commitments;
- unresolved contradictions.

12.2 Instrument 2: Factorial scenario system

Systematically vary:

- target identity;
- relationship;
- in-group versus out-group;
- harm severity;
- intentionality;
- foreseeability;
- legality;
- authority instruction;
- audience;
- scarcity;
- threat;
- personal cost;
- group benefit;
- certainty of evidence;
- power asymmetry;
- role obligation;
- reversibility.

Primary output:

Context coefficients indicating which features change judgment, intention, punishment, tolerance, or willingness to act.

12.3 Instrument 3: Trade-off battery

Candidate conflicts include:

- welfare versus loyalty;
- autonomy versus protection;
- equality versus hierarchy;
- truth versus relationship preservation;
- individual rights versus collective security;
- rule consistency versus exceptional need;
- immediate harm versus uncertain future benefit;
- mercy versus deterrence;
- personal integrity versus institutional obedience;
- in-group protection versus universal scope.

Primary output:

Priority thresholds and conflict-resolution rules.

12.4 Instrument 4: Public-private divergence protocol

Compare:

- anonymous response;
- identified response;
- in-group audience;
- out-group audience;
- authority-observed response;
- expected behavior;
- intended behavior;
- reported past behavior;
- behavior under opportunity.

Primary output:

A divergence profile that distinguishes strategic presentation, conformity, concealment, uncertainty, and genuine revision.

12.5 Instrument 5: Ecological momentary assessment

Brief repeated prompts capture:

- what happened;
- what the person believed was happening;
- what mattered;
- what they felt required or permitted to do;
- what they wanted;

- what they intended;
- what they did;
- what constrained them;
- what followed;
- whether their judgment changed.

Primary output:

Within-person variation and real-world base rates.

12.6 Instrument 6: Developmental narrative interview

Trace:

- family rules and discipline;
- early authority relationships;
- religious and cultural learning;
- peer influence;
- experiences of injustice;
- significant transgressions;
- punishment and repair;
- identity transitions;
- ideological conversions;
- disillusionment;
- value conflicts;
- moral injury;
- moments of revision.

Primary output:

A developmental provenance map rather than a deterministic causal claim.

12.7 Instrument 7: Epistemic update protocol

Vary:

- source credibility;
- evidence amount;
- consensus;
- ambiguity;
- identity relevance;
- emotional threat;
- group endorsement;
- disconfirming information.

Primary output:

Belief calibration, source trust, identity-protective resistance, and revisability.

13. Relationship to personality and other Tessera systems

13.1 Personality

Personality is a modulator of expression, not normative content.

Examples:

- conscientiousness may increase consistency with an accepted rule;
- agreeableness may increase relational accommodation;
- threat sensitivity may narrow scope under danger;
- openness may increase exposure to alternative frameworks;
- assertiveness may affect willingness to act publicly;
- emotional volatility may alter behavior under conflict.

None of these traits specifies which ideology, beneficiary, principle, or moral target the person will endorse.

```
Normative expression = f(  
  normative content,  
  personality,  
  affect,  
  cognition,  
  identity,  
  capability,  
  ecology,  
  constraint,  
  time  
)
```

13.2 Reconciliation

Reconciliation should examine apparent mismatches such as:

- declared value versus behavior;
- principle versus exception;
- private belief versus public response;
- intention versus capability;
- intended benefit versus actual harm;
- empathy versus restricted scope;
- sophisticated reasoning versus biased premises;
- identity commitment versus new evidence.

Reconciliation does not automatically excuse harm or invalidate the participant's report. It organizes competing explanations for review.

13.3 Karma and Ethical Integrity Matrix

The existing 12-dimensional Ethical Integrity and Karma material should be retained as a provisional second-layer overlay.

Existing dimension	Descriptive redevelopment within TNOA
Authentic Presentation	Public-private consistency and strategic presentation
Relational Equity	Distributional rule, entitlement, autonomy, and exploitation permission
Self-Regard Calibration	Self-priority, status claim, and entitlement structure
Material Orientation	Resource and status prioritization
Affective Empathy	Detection of and affective response to others' states
Accountability Orientation	Responsibility acceptance and attribution pattern
Intentional Orientation	Intended beneficiary, intended burden, and accepted harm
Moral Reasoning	Justificatory form, scope, and complexity
Constraint Internalization	Authority locus and degree of internalization
Moral Sensitivity	Consequence detection and concern weighting
Moral Courage	Cost-bearing consistency under pressure
Development Orientation	Revisability and response to contradiction

These dimensions should not be automatically summed. A person may display high commitment, courage, empathy, or reasoning complexity in service of harmful aims. The profile and its contradictions are the finding.

13.4 Object, cognitive, identity, and behavioral architecture

TNOA fits the Tessera sequence:

```
Object Architecture
-> Cognitive Architecture
-> Identity Architecture
-> Normative Orientation Architecture
-> Behavioral Architecture
-> Consequence and Developmental Update
```

This ordering is functional rather than permanently fixed. Normative orientation also feeds backward into perception, memory, identity, and future object meaning.

14. Scoring and representation policy

14.1 No default global score

The initial model should produce profiles, networks, thresholds, and contextual coefficients. A global morality, integrity, or normative-quality score is prohibited unless a later validation program establishes a clear construct and declared benchmark.

14.2 Appropriate score types

Depending on the construct, acceptable representations may include:

- categorical status;
- continuous estimate;
- rank order;
- probability distribution;
- threshold;

- network centrality;
- context coefficient;
- within-person variability;
- source-confidence rating;
- contradiction flag;
- qualitative code;
- longitudinal trajectory.

14.3 Reflective versus formative constructs

Internal consistency metrics are appropriate only when items are expected to reflect one latent variable. Value profiles, ideology networks, episode configurations, and contextual composites may be formative or network-based and should not be rejected merely because they do not behave like a unidimensional scale.

14.4 Inference confidence

Suggested confidence labels:

- Directly reported;
- Behaviorally observed;
- Repeatedly corroborated;
- Context-limited;
- Inferred with moderate support;
- Speculative hypothesis;
- Contradicted;
- Insufficient evidence.

15. Validation roadmap

Phase 0: Construct governance

Target: Months 0-3

Deliverables:

- Normative Orientation Charter;
- descriptive-evaluative boundary policy;
- master glossary;
- evidence classification rules;
- privacy and ethics protocol;
- cultural review framework;
- inference-confidence rules.

Gate:

No item or score development begins until construct boundaries are reviewed.

Phase 1: Independent domain analysis

Target: Months 3-9

Each workstream produces:

- Construct Development Packet;
- literature map;
- facet proposal;
- developmental account;
- ecological account;
- measurement candidates;
- discriminant-validity targets;
- graph specification.

Gate:

Every construct must be distinguishable from its closest neighbors.

Phase 2: Semantic ontology and coding system

Target: Months 6-12

Deliverables:

- Normative Decision Episode schema;
- controlled vocabulary;
- coding manual;
- context taxonomy;
- authority taxonomy;
- justification taxonomy;
- harm and beneficiary taxonomy;
- uncertainty and provenance system;
- graph prototype.

Gate:

Independent coders must represent the same episode with acceptable agreement without being forced into moral evaluation.

Phase 3: Measurement prototype

Target: Months 9-18

Deliverables:

- interview protocol;
- scenario library;
- trade-off tasks;

- public-private protocol;
- developmental narrative interview;
- ecological momentary assessment;
- epistemic update tasks;
- preliminary coding and scoring rules.

Gate:

Cognitive interviews should show that participants understand the measures across relevant education levels, ages, cultures, and ideological positions.

Phase 4: Construct validation

Target: Months 18-30

Evaluate:

- inter-rater reliability;
- test-retest reliability where stability is expected;
- within-person variability where context sensitivity is expected;
- convergent validity;
- discriminant validity;
- measurement invariance;
- known-groups comparisons;
- social-desirability sensitivity;
- response-process validity;
- criterion validity where ethically appropriate.

Gate:

No holistic profile is produced until the independent domains demonstrate adequate measurement.

Phase 5: Process validation

Target: Months 24-36

Test questions such as:

- Does value activation alter deontic judgment?
- Does perceived threat narrow normative scope?
- Does public observation increase conformity?
- Does authority instruction alter responsibility attribution?
- Does high constraint load reduce reflective override?
- Do changed factual premises alter judgment without changing values?
- Does perceived control alter behavior while leaving intention intact?
- Do consequences update values, principles, identity, or merely strategy?

Gate:

Causal language is restricted to processes supported by design and evidence.

Phase 6: Bioecological and longitudinal validation

Target: Years 3-5

Track:

- family and peer socialization;
- school and occupational systems;
- religious and political institutions;
- media environments;
- social mobility;
- trauma and threat;
- group membership;
- historical events;
- identity transitions;
- ideological stability and revision.

Analyses should distinguish:

- between-person differences;
- within-person change;
- developmental trajectories;
- period effects;
- cohort effects;
- person-context feedback loops.

Gate:

The model must represent both stability and change without treating all variation as inconsistency.

Phase 7: Integration and translation

Target: Years 5-7

Deliverables:

- holistic Tessera integration;
- reporting standard;
- optional ethical benchmark overlays;
- research API and graph interface;
- training manual;
- bias audit;
- cross-cultural adaptation protocol;
- use restrictions;
- versioned evidence ledger.

16. Holistic profile output

A completed profile should contain seven sections.

16.1 Normative content profile

- values;
- principles;
- norms;
- ideologies;
- moralized domains;
- ethical reasoning modes;
- factual premises.

16.2 Normative structure profile

- scope;
- authority;
- justification;
- responsibility;
- harm tolerance;
- beneficiary orientation;
- punishment and repair model.

16.3 Contextual activation profile

- roles;
- relationships;
- audiences;
- power;
- threat;
- scarcity;
- institutional setting;
- group identity;
- public-private conditions.

16.4 Intention-enactment profile

- motives;
- intentions;
- plans;

- control;
- opportunity;
- constraints;
- behavior;
- outcomes.

16.5 Developmental provenance profile

- learning sources;
- formative relationships;
- cultural systems;
- significant events;
- reinforcement history;
- narrative integration;
- historical changes.

16.6 Coherence and flexibility profile

- internal contradictions;
- contextual exceptions;
- public-private divergence;
- self-other asymmetry;
- value-principle alignment;
- judgment-behavior divergence;
- revisability;
- identity-protective resistance.

16.7 Evidence and uncertainty profile

Every conclusion should state:

- data source;
- direct report or inference;
- number of supporting observations;
- context coverage;
- confidence;
- alternative explanations;
- evidence classification;
- validation status.

17. Development sequence and dependencies

Recommended build order:

- 1 Values.
- 2 Principles.
- 3 Norms.
- 4 Ideology and factual premises.
- 5 Morality and moralization.
- 6 Ethics and justificatory repertoire.
- 7 Motives.
- 8 Intentions.
- 9 Behavior and consequences.
- 10 Contextual modulation.
- 11 Developmental and narrative integration.
- 12 Optional ethical evaluation overlays.

Dependency logic:

- Values precede principles because principles operationalize values.
- Principles and norms must be separated before ideology can be mapped cleanly.
- Factual premises must be separated from values before disagreement is interpreted.
- Motives and intentions must be separated before behavior is used as evidence.
- Context must be measured at every stage rather than appended afterward.
- Ethical evaluation comes last so that descriptive measurement is not bent toward predetermined conclusions.

18. Repository implementation plan

18.1 Documentation

The repository should maintain:

- this Markdown specification as the editable source;
- a matching PDF as a review and distribution artifact;
- README links to both;
- manifest classification;
- a generated repository tree;
- validation workflow checks for paired specification files and governance references.

18.2 Proposed future documentation set

```
docs/normative-orientation/  
  glossary.md  
  construct-governance.md  
  workstreams/  
    values.md  
    principles.md  
    norms.md  
    ideology.md  
    morality.md  
    ethics.md  
    epistemic-premises.md  
    motives.md  
    intentions.md  
    behavior-and-consequences.md  
  measurement/  
    interview-protocol.md  
    scenario-design.md  
    tradeoff-battery.md  
    public-private-protocol.md  
    ecological-momentary-assessment.md  
    developmental-narrative-interview.md  
  ontology/  
    normative-decision-episode.md  
    graph-schema.md  
    controlled-vocabulary.md  
  validation/  
    validation-roadmap.md  
    evidence-ledger.md
```

This proposed tree is a development scaffold, not a requirement to create all files immediately.

18.3 Versioning

Suggested specification versions:

- 0.x: conceptual and prototype development;
- 1.0: construct boundaries and semantic schema approved for formal pilot;
- 2.0: initial validation completed for core workstreams;
- 3.0: longitudinal and cross-cultural evidence sufficient for an integrated reporting standard.

18.4 Change control

Any future change that alters construct boundaries, scoring, evaluation, or data interpretation should include:

- overwrite or reinterpretation notice;
- affected files;
- evidence rationale;
- migration impact;
- backward-compatibility note;
- validation consequences;
- approval status.

19. Risk register

Risk	Consequence	Primary safeguard
Moral evaluation leaks into description	Biased profiles and circular scoring	Keep descriptive, explanatory, and evaluative layers separate.
Personality is treated as moral content	Trait-based moral stereotyping	Personality may modulate expression but cannot determine normative content.
Behavior is treated as transparent intent	False attribution	Separate motive, intention, control, opportunity, action, and consequence.
Context becomes an afterthought	Low ecological validity	Use episode-level and longitudinal person-in-context measurement.
Articulate reasoning is mistaken for ethical quality	Rewards verbal skill or rationalization	Code justificatory form separately from scope, premises, beneficiaries, and harm.
Internal consistency is overused	Formative or network constructs are rejected	Match psychometrics to construct structure.
Automated inference exceeds evidence	Harmful certainty	Preserve provenance, confidence, alternatives, and human review.
Cultural variation is mistaken for deficiency	Biased interpretation	Use emic elicitation, invariance testing, and cultural review.
Harmful ideology is described without critical boundary	Apparent normalization	Descriptive representation does not equal endorsement; evaluative benchmarks remain explicit.
Sensitive identity data is misused	Privacy or discrimination harm	Data minimization, access control, purpose limitation, deletion policy, and use restrictions.

20. Acceptance criteria for Version 0.1

Version 0.1 is complete when:

- the umbrella definition is explicit;
- all subordinate constructs are separately defined;
- the layered architecture is documented;
- the bioecological organization is specified;
- every workstream has required distinctions, methods, and outputs;
- the Construct Development Packet is reusable;
- the Normative Decision Episode schema is specified;
- graph nodes and relations are defined;
- the multimethod battery is outlined;
- the relationship to personality, reconciliation, Karma, and Tessera engines is explicit;
- scoring restraints are documented;
- validation phases and gates are defined;
- repository governance and future file structure are specified;
- uncertainty and evidence status remain visible.

Appendix A. Construct Development Packet template

Construct name:
Document version:
Status:
Evidence level:
Owner or workstream:
Date reviewed:

1. Formal definition
2. Plain-language definition
3. Necessary features
4. Optional features
5. Exclusion criteria
6. Neighboring constructs
7. Common conflations
8. Construct type
9. Proposed facets
10. Reflective/formative/network status
11. Developmental account
12. Bioecological account
13. Mechanistic account
14. Measurement modes
15. Scoring/coding rules
16. Missing and uncertain data
17. Validity risks
18. Cross-cultural considerations
19. Tessera graph implementation
20. Evidence provenance
21. Advancement gate decision
22. Open questions

Appendix B. Normative Decision Episode example

Context: Workplace team meeting
Role: Project lead
Audience: Direct reports and supervisor
Target: Team member who missed a deadline
Perceived situation: Repeated failure threatens the project
Available actions: Private inquiry, public criticism, reassignment, formal discipline
Normative status: Public criticism judged permissible but not required
Activated values: Efficiency, accountability, fairness
Applied principle: People should answer for preventable failures
Perceived norm: Strong leaders address failure immediately
Ideological frame: Merit and responsibility
Factual premise: The delay was believed to be controllable
Motive: Restore control and signal standards
Intended beneficiary: Project and wider team
Foreseen harm: Embarrassment and reduced trust
Accepted harm: Moderate embarrassment judged necessary
Intention: Confront publicly
Perceived control: High
Behavior: Public criticism
Actual outcome: Compliance increased; trust decreased
Post-event revision: Later evidence showed the delay was caused by an unreported family emergency
Revisability question: Does the accountability principle change, or only the factual premise and chosen means?

Appendix C. Initial research questions

- 1 Can values, principles, norms, and ideology be measured with adequate discriminant validity?
- 2 Which contextual manipulations most reliably reveal normative scope and exception rules?
- 3 How stable are deontic judgments within persons across role, audience, threat, and target changes?

- 4 Can factual-premise updates change judgment without changing value hierarchy?
- 5 Which profile features predict intention-behavior divergence?
- 6 Does constraint load narrow justificatory repertoire or only behavioral options?
- 7 How do public-private differences vary across culture and power relations?
- 8 Can longitudinal data distinguish genuine revision from strategic re-description?
- 9 Which parts of the existing Karma matrix function as descriptive axes, and which require declared ethical benchmarks?
- 10 What minimum data are required before the system may infer a stable identity-level normative orientation?

Appendix D. Current status statement

TNOA is presently a **provisional integrative specification**. It provides a development map, semantic ontology, measurement plan, and validation roadmap. It is not yet an established psychological theory, validated assessment, diagnostic instrument, or moral-authority system.